

Visualization - Intro (Lecture 2)

Peter Ganong and Maggie Shi

January 12, 2026

Citing our sources

- Schwabish: “Better Data Visualizations” ([link to purchase](#))
- Healy: “Data Visualization” ([link to full text](#))
- Heer, Moritz, VanderPlas, and Craft `altair` textbook: ([link to full text](#))

Why visualize?

Why visualize?: roadmap

- Introduce **Anscombe's Quartet** – a classic data visualization example
- Look at the data directly, calculate summary statistics, then plot
- Discuss what observations can be made at each step

Anscombe's Quartet

- Introduced by statistician **Francis Anscombe** in 1973
- Consists of **four small datasets** – each with two variables: **x** and **y**

The raw data

I		II		III		IV	
x	y	x	y	x	y	x	y
4	4.3	4	3.1	4	5.4	8	5.3
5	5.7	5	4.7	5	5.7	8	5.8
6	7.2	6	6.1	6	6.1	8	6.6
7	4.8	7	7.3	7	6.4	8	6.9
8	7.0	8	8.1	8	6.8	8	7.0
9	8.8	9	8.8	9	7.1	8	7.7
10	8.0	10	9.1	10	7.5	8	7.9
11	8.3	11	9.3	11	7.8	8	8.5
12	10.8	12	9.1	12	8.2	8	8.8
13	7.6	13	8.7	13	12.7	8	12.5
14	10.0	14	8.1	14	8.8	19	5.6


Summary statistics

Let's also compute some summary statistics

	series	Mean X	Mean Y	Var(x)	Var(y)	Corr(x,y)	Regression
0	I	9.0	7.5	11.0	4.1	0.8	$y = 0.50x + 3.0$
1	II	9.0	7.5	11.0	4.1	0.8	$y = 0.50x + 3.0$
2	III	9.0	7.5	11.0	4.1	0.8	$y = 0.50x + 3.0$
3	IV	9.0	7.5	11.0	4.1	0.8	$y = 0.50x + 3.0$


- All four datasets have basically **identical** summary statistics and regression coefficients!
- But are they really capturing the same relationships between x and y?

Anscombe's quartet, visualized




Discussion question: what observations can we now make about the similarities and differences across the series?

Observations from the visualization



- Series I, II, and III illustrate very different positive relationships between x and y
 - I: noisy and approximately linear
 - II: quadratic and fully deterministic
 - III: very linear, with exception of an outlier

Observations from the visualization



- The two outliers in III and IV stand out much more quickly
- The positive relationship suggested by the regression in series IV is just an artifact of the outlier

Why visualize?: summary

- Anscombe's Quartet illustrates that while summary statistics are useful, we can't rely on them alone
- The best way to detect patterns is to visualize your data

Introduction to Vega- Lite and `altair`

Roadmap

- What's different from `matplotlib`?
- What is Vega-Lite?
- What is `altair`?
- First plot – image and then grammar

matplotlib is imperative

```
1 import matplotlib.pyplot as plt
2
3 plt.ylabel("Y axis")
4 plt.figure(figsize=(6, 4))
5 plt.title("A Line Plot")
6 plt.plot([1, 2, 3], [4, 1, 6], color='red', marker='o')
7 plt.xlabel("X axis")
8
9 plt.grid(True)
10 plt.show()
```

- **Imperative:** you tell computer directly how to draw graph
- Each graphic element is layered on one-by-one, but not organized
- Not clear what type of plot this is
- Not clear how data is being plotted (which axis is which?)

Declarative approaches to visualization

- Good graphics packages are **declarative**: you provide a high-level specification of *what* you want in the visualization.
- Importantly, you do so in an **organized** way.
- Declarative visualization approaches have three inputs:
 - **Data**
 - **Graphical marks**: the “type” of plot – bar, scatter, etc.
 - **Encoding channels**: x-axis, y-axis, colors, etc


We are going to use *Vega-Lite* + `altair` in this class.

What is Vega-Lite and Altair?

- Vega-Lite is a “grammar” of interactive graphics
 - Under the hood: [JSON \(JavaScript Object Notation\)](#) is used to record Vega-Lite specifications
 - *Note: you will not have to write Vega-Lite manually*
- [altair](#) is a Python package allows Python to write Vega-Lite

Example: Graph vs. Grammar

Plot:



Vega-lite (JSON) underlying the plot:

```
1 {  
2   "data": {"url": "data/seattle-weather.csv"},  
3   "mark": "bar",  
4   "encoding": {  
5     "x": {"timeUnit": "month", "field": "date"}  
6     "y": {"aggregate": "mean", "field": "precipitation"}  
7   }  
8 }
```

- This is the code that is “under the hood” of the graphic – *not* the Python code you write.
- Discussion question: can you tell what each line of text means?

Source.

Why the emphasis on grammar?

- Packages for making graphics and coding languages change over time
- We chose to teach a package with an underlying grammar because we are trying to foreground the *conceptual* aspects of data visualization
- Insights are portable, even as the particular package/language you use changes over time

Summary

- Use a declarative approach
- Grammar: Vega-lite
- Python package to write Vega-lite: `altair`
- Gives a coherent conceptual representation underlying a plot

Introduction to altair

What is `altair`? + roadmap

`altair` is a Python API (Application Programming Interface) that generates Vega-Lite specifications.

Roadmap:

- Load package
- Define dataset we'll work with through rest of class

Imports and Renderer

```
1 import pandas as pd  
2 import altair as alt
```

Note: depending on your environment, you may also need to specify a renderer for `altair` (see [Leja's Ed post](#) as one example)

- If you run into this, please read the documentation for [Displaying Altair Charts](#).
- If that fails, post in Ed and bring your question to lab.

Tidy data: weather data

- Visualization in `altair` begins with “tidy” data frames
 - Each variable is a column
 - Each observation is a row
 - Each value is a cell
- A simple data frame (`df`) containing the average precipitation (`precip`) for a given `city` and `month` :

```
1 df = pd.DataFrame({  
2     'city': ['Seattle', 'Seattle', 'Seattle', 'Seattle', 'New York', 'New Y  
3     'month': [1, 4, 8, 12, 1, 4, 8, 12, 1, 4, 8, 12],  
4     'precip': [3.12, 2.68, 0.87, 5.31, 2.1, 3.94, 4.13, 3.58, 3.3, 3.62, 3.  
5 })
```

Tidy data: weather data

```
1 df.head()
```

	city	month	precip
0	Seattle	1	3.12
1	Seattle	4	2.68
2	Seattle	8	0.87
3	Seattle	12	5.31
4	New York	1	2.10

Aside: in-class examples and tidy data

- Data you will work with is most often *not* as tidy as the in-class examples
- Much of the work you will do as an analyst is kind of this “unglamorous” work!
- Using `altair` to make the plot is often the last step
 - In lecture, we’ll skip the data cleaning to focus on ideas and skills about how to do visualization well
 - In problem sets + project, you will start with messy data that will require cleaning
 - Why don’t we cover data cleaning in lecture? Every dataset needs different steps for cleaning; best way to teach it is therefore through hands-on exercises in the problem sets.

Summary

- `altair` is an API that enables Python to “speak” in Vega-Lite’s grammar
- Input to `altair`: tidy data

Building a first chart

Building a first chart: roadmap

- Incrementally build our first chart in `altair`
- Then build our first aggregated chart

altair ingredients

Recall that we need three inputs to a *Vega-lite* chart:

- **Data**
- **Graphical marks:** the “type” of plot - bar, scatter, etc.
- **Encoding channels:** x-axis, y-axis, colors, etc

Data: the **Chart** object

```
1 mychart = alt.Chart(df)
```

- We have defined the **mychart** object and passed it the data
- But nothing has been plotted yet – still need mark and encoding

Mark: point


```
1 mychart = alt.Chart(df).mark_point()  
2  
3 mychart
```



- Now, declare the mark: a `mark_point()`
- We haven't declared what's on the axes
- So actually this is *all the data points*, located in the same place

Encoding: one point per city on y-axis


```
1 mychart = alt.Chart(df).mark_point().encode(  
2     alt.Y('city')  
3 )  
4  
5 mychart
```



- Using `alt.Y`, we've separated data by one attribute along the y-axis: city
- We haven't encoded anything on the x-axis yet!
- So underneath each point, we have multiple points overlapping within each city

Encoding: xy coordinates


```
1 mychart = alt.Chart(df).mark_point().encode(  
2     alt.X('precip'),  
3     alt.Y('city')  
4 )  
5  
6 mychart
```



- Process: Code is super-duper readable.
- Note that we've directly plotted all the data
- Substance: *Seattle exhibits both the least-rainiest and most-rainiest months!*

Quick Aggregation


```
1 mychart = alt.Chart(df).mark_point().encode(  
2     alt.X('average(precip)'),  
3     alt.Y('city'))  
4 )  
5  
6 mychart
```



- Say instead of plotting all data points, we want to transform and aggregate it first
- `altair` has a few “shorthand” aggregation functions. like `average()`
- We will discuss more complicated aggregation in week 3

Changing Marks is Straightforward

```
1 mychart = alt.Chart(df).mark_bar().encode(  
2     alt.X('average(precip)'),  
3     alt.Y('city'))  
4 )  
5  
6 mychart
```



Syntax: Understanding Altair's shorthands.


Three ways to say the same idea

```
1 # what we will continue to use
2 alt.X('average(precip)')
3
4 # shorter
5 x = 'average(precip)'
6
7 # longer
8 alt.X(aggregate='average', field='precip', type='quantitative')
```

- Going forward, we will primarily use the first
- But you may see all three in problem sets, lab assessments, etc.

Customizing a plot: colors and labels

```
1 alt.Chart(df).mark_point(color='firebrick').encode(  
2     alt.X('precip', title='Rain (in)'),  
3     alt.Y('city', title='City'),  
4 )
```




- Add customizations and labels where you declare the relevant mark or encoding
- This is what we mean by “grammar” – code is organized in a consistent, readable way

Comparison to `matplotlib`


altair

```
1 mychart = alt.Chart(df).mark_point(color='firebrick').encode()
2     alt.X('average(precip)', title='Rain (in)'),
3     alt.Y('city', title='City'),
4 )
5
6 mychart
```



matplotlib

```
1 import matplotlib.pyplot as plt
2 avg_precip = df.groupby('city', as_index=False)[['precip']].mean()
3
4 plt.figure(figsize=(4, 3))
5 plt.xlabel('Rain (in)')
6 plt.scatter(avg_precip['precip'], avg_precip['city'], color='red')
7 plt.grid(True, axis='x')
8 plt.ylabel('City')
9 plt.show()
```



- **matplotlib**: have to first make a separate, collapsed **pandas** dataframe
- x and y labels defined separately from code for plot

Visualization guidelines

- All axes and units are properly labeled and legible
- No words or data points are cut off in your final output
- Encodings should be sensible/appropriate – *more in this next lecture*

Building a first chart: summary

- Everything begins with a `Chart(data)`
- Every `Chart` needs a `mark`
- Every `Chart` needs guidance how to encode the data in terms of `marks`
- Simple chart formatting:
`mark_point(color='firebrick'), alt.X(title = '...')`

Do-pair-share

- *Do* – make a plot on your own
 - *Pair* – compare your results with person next to you
 - *Share* – discuss results as a class
1. Open `viz_1_intro/viz_1_dps.qmd` in VSCode
 2. If you have the Quarto and Jupyter extensions installed + dap conda environment set up ([link](#)), you should be able to directly “Run Cell” within
 - `.qmd` file

Data Transformation: Do-pair-share

- Make a bar plot showing the **lowest** rainfall for each city in the dataset.
- Starter code in `viz_1_intro/viz_1_dps.qmd` file in student repo:

```
1 import pandas as pd
2 import altair as alt
3 df = pd.DataFrame({
4     'city': ['Seattle', 'Seattle', 'Seattle', 'Seattle', 'New York', 'New York', 'New York', 'N
5     'month': [1, 4, 8, 12, 1, 4, 8, 12, 1, 4, 8, 12],
6     'precip': [3.12, 2.68, 0.87, 5.31, 2.1, 3.94, 4.13, 3.58, 3.3, 3.62, 3.98, 2.56]
7 })
```

- Hint: Altair's aggregation methods are [here](#)


Multiple Views

Multiple Views: roadmap

- Introduce `mark_line()`
- Multiple marks
- Multiple panels

mark_line()

```
1 line = alt.Chart(df).mark_line().encode(  
2     alt.X('month', title = "Month"),  
3     alt.Y('average(precip)', title = "Average rainfall (in.)")  
4 )  
5  
6 line
```




Multiple Marks

- Now say that we want to add layer scatter points on top of our line to highlight the unit of observation
- `altair` grammar allows you to “layer” one set of marks on another with intuitive syntax: `+`

Multiple Marks

```
1 line = alt.Chart(df).mark_line().encode(
2     alt.X('month', title = "Month"),
3     alt.Y('average(precip)', title = "Average rainfall (in.)")
4 )
5 point = alt.Chart(df).mark_circle(color='firebrick').encode(
6     alt.X('month', title = "Month"),
7     alt.Y('average(precip)', title = "Average rainfall (in.)")
8 )
9 line + point
```





Multiple Panels

- Now say we want to place another plot *next to* our original plot
- Again, `altair` grammar allows for fairly intuitive syntax: |

Multiple Panels



```
1 line_min = alt.Chart(df).mark_line().encode(
2     alt.X('month', title = "Month"),
3     alt.Y('min(precip)', title = "Lowest rainfall (in.)")
4 )
5
6 line_avg | line_min
```



Discussion: while technically correct, there's something not quite right about these charts...

Multiple Panels, take 2

```
1 line_min = alt.Chart(df).mark_line().encode(
2     alt.X('month', title = "Month"),
3     alt.Y('min(precip)', title = "Lowest rainfall (in.)")
4 )
5 combined = line_avg | line_min
6 combined = combined.resolve_scale(y='shared')
7 combined
```





`resolve_scale()` is applied to `combined`, not `line_avg` or `line_min`

Organized and easy to understand!

Equivalent graph in `matplotlib`

Of course, you can make a very similar plot in `matplotlib`



Equivalent code in matplotlib

```
1 avg_by_month = df.groupby('month', as_index=False)['precip'].mean()
2 fig, axs = plt.subplots(1, 2, figsize=(12, 4), sharey=True)
3 axs[1].set_ylabel('Min of precip')
4 axs[0].plot(avg_by_month['month'], avg_by_month['precip'])
5 min_by_month = df.groupby('month', as_index=False)['precip'].min()
6 axs[1].scatter(min_by_month['month'], min_by_month['precip'])
7 axs[0].scatter(avg_by_month['month'], avg_by_month['precip'])
8 axs[0].set_ylabel('Average of precip')
9 axs[0].set_xlabel('Month')
10 axs[1].plot(min_by_month['month'], min_by_month['precip'])
11 axs[1].set_xlabel('Month')
12
13 plt.tight_layout()
14 plt.show()
```

Multiple Views – summary

- New syntax:
 - `plot1 + plot2` for multiple marks on same panel
 - `plot1 | plot2` for multiple panels
- `altair` allows you to combine multiple plots in an organized and intuitive way
- Usefulness of grammar and declarative approach becomes more apparent when extending beyond a single plot

